

Securing Windows Subsystem for Linux A Behavioral Detection Approach

1st Andrei Mermeze

dept. name of organization (of Aff.)

name of organization (of Aff.)

Cluj-Napoca, Romania

andrei.mermeze@gmail.com

2nd Phd. Radu Dragos

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address

Abstract—The release of Windows Subsystem for Linux (WSL) revealed a whole new attack surface, comprised of kernel drivers as well as user-mode services. The performed research is going to reveal how behavioral detection techniques can be applied for detecting potential threats (bashware) that abuse WSL. This paper will provide an insight into the security issues created by this subsystem, while also explore both Kernel-Mode and User-Mode based detection heuristics and techniques in order to identify and block this new type of malware. First sections focus on WSL internals, whereas the next sections present what mechanisms can be used to acquire the needed information for identifying bashware, and finally some heuristic behavioural based approaches on identifying the bashware.

Index Terms—Windows Subsystem for Linux Security Behavioral Analysis Malware Bashware Detection Event Tracing Minifilter Monitoring Kernel Driver

I. INTRODUCTION

Windows Subsystem for Linux was first released in the anniversary update and it provides a way of executing Linux ELF 64 bit binaries on native Windows 10. Since WSL is not a virtualization based system, the Linux processes running in it can access any resource on the computer, making it a dangerous threat until a AV Solutions are updated to support this new type of processes.

II. OVERVIEW OF WSL

Windows Subsystem for Linux is a new and complex component in Windows that implements a lot of new functionality (in syscalls), parsing (ELF files parsing) and is based existing Windows technologies, like pico providers and pico processes.

A. Minimal Processes and Pico Processes

A minimal process has a parent, protection level, name and security token. It has no initial thread, no process environment block (PEB), no ntdll, and its address space is empty. Its threads are called "minimal threads", and, similarly to minimal processes, they have no thread environment block (TEB). A pico process is a minimal process, but it has an associated pico provider, which handles the system calls, exceptions, pico process or threads creation and termination.

B. Pico Providers

A pico provider is a kernel driver that implements the required functionalities to handle pico process events. Registering as a pico provider is done by calling the PsRegisterPicoProvider API, however, calling this API requires that the PspPicoRegistrationDisabled is set to FALSE. This value is set to false before any other third party driver is loaded, therefore, at least for now, it is not possible to implement custom pico providers. In order to secure pico providers, they also register with PathGuard in order to protect its syscalls handlers. Kernel Patch Protection, also known as PatchGuard, was designed to protect kernel structures from being patched. This renders the classic linux monitoring solution (syscall hooking) useless, as any attempt to hook the syscalls will result in a BSOD (Blue Screen of Death).

Lxcore.sys does the pico provider registration in the LxInitialize function, which is called by lxss.sys in its DriverEntry.

```
result = PsRegisterPicoProvider(&v4, &LxpRoutines);
```

Fig. 1. lxcore.sys pico provider registration

C. Syscalls

Currently, as of Windows 10 1709, 242 Linux syscalls are implemented in lxcore.sys. These syscalls are implemented as callbacks that reside in a table (lxcore!LxpSyscalls), and are called from lxcore's LxpSysDispatch callback, which was registered with PsRegisterPicoProvider.

D. File System

- VFS
- VolFs - /mnt/c
- DrvFs - /drv
- SysFs - /sys
- TmpFs - /tmp
- ProcFs - /proc

III. EXPLORING THE ATTACK SURFACE

Lxcore is the core kernel-mode component of WSL that implements the Linux syscalls, multiple filesystems, and all the logic for running linux binaries in Windows. The driver, being responsible with providing the kernel support for user processes, implements many parsing algorithms (i.e. ELF header parsing, user strings parsing), making it prone to overflows, off-by-ones and many other issues that can be used in developing exploits, as we will see in the next section.

IV. KNOWN VULNERABILITIES

A. *Execve Exploit*

Was discovered by Saar Amar, and was given the following CVE: CVE-2018-0743. The exploit leverages an integer overflow in `Lxcore!LxpUtilReadUserStringSet` in order to run privileged code. The shellcode elevates a process given by its pid. In this paper we are going to explain in detail the algorithm we came up with.

B. *Local Denial of Service*

While experimenting with Event Tracing for Windows and reversing Lxcore, we've discovered an issue in Lxcore allowing us to trigger a BSOD from an unprivileged user-mode process. However, due to responsibility disclosure, we will not go into any more detail about this issue until it is patched by Microsoft.

V. MONITORING WSL ACTIVITY

There are multiple ways of monitoring WSL activity, which, if correctly combined, can provide enough information in order to identify potentially malicious activity.

- User-Mode Hooking Framework - API usage monitoring
- File System Minifilter - file system I/O and process monitoring
- Event Tracing for Windows
- Windows Filtering Platform - network monitoring

While we've studied and experimented with all of the above, we focused only on the file system minifilter approach, therefore we won't go into great detail about the other methods.

A. *User-Mode hooking driven WSL monitoring*

In order to obtain more granular monitoring capabilities, a function hooking framework is required. The ability of synchronously monitoring specific API usage opens up a range of possibilities, from exploit detection to syscall graph based detection heuristics. In order to load a shared library into any Linux pico process, an entry with the path to the library must be added into the `ld.so.conf`. Since any process with root privileges (inside the linux system) can update the said config file, we must protect it from a kernel driver, so

that we are sure that our shared library isn't removed. This can be done by filtering `IRP_MJ_WRITE` operations on that file and readding the entry if removed.

When loaded in a linux process, the library hooks the functions we want to monitor (i.e. `ptrace`, `open`) and synchronously notifies a service via local sockets about every hooked function call, informing the service about exploitation intent and failing the function call if needed.

We consider this method unreliable against smarter bashware that actively checks and avoids hooked functions, and possibly completely useless starting with Windows RS5 if executable memory can't be allocated anymore.

B. *Using A File System Minifilter to Monitor WSL Activity*

This is the most reliable way of monitoring and blocking potentially malicious Linux applications. Unlike a user-mode hooking framework approach, using a Windows driver provides mechanisms of monitoring that cannot be bypassed or tampered with by the monitored processes, making it a very reliable and secure approach.

The driver we've implemented filters file system I/O and also keeps track of active processes. In the next sections we are going to present how we implemented the driver.

1) File System Filtering: In Windows I/O request packets (IRPs) are used to communicate between drivers. For example, in the context of the file system, IRPs are used to describe file I/O operations, like file read (`IRP_MJ_READ`) or file create (`IRP_MJ_CREATE`). A minifilter driver can register pre and post callbacks for all IRPs by calling `FltRegisterFilter`. As the name suggests, the pre callback is called before the operation is completed and the post callback is called after the operation is completed. For example, the pre callback for an `IRP_MJ_CREATE` would show the intent of creating or opening a file, but only in the post callback we would know if it succeeded or not (file might not exist in case of an open).

We've used this technology in order to monitor linux processes as well as win32 processes file system activity. We have decided to only monitor win32 processes that are part of a linux process tree (an ancestor of the process is a linux process) in order to reduce the performance overhead.

When filtering, we are mostly interested in 2 fields of the `FLT_CALLBACK_DATA` structure, while the `FILE_OBJECT` is taken from the `FLT_RELATED_OBJECTS` structure.

```
0: kd> dt fltmgr!_FLT_CALLBACK_DATA
...
0x10 Iopb          : Ptr64 _FLT_IO_PARAMETER_BLOCK
...
```

```

0x50 RequestorMode: Char

0: kd> dt fltmgr!_FLT_RELATED_OBJECTS
...
0x20 FileObject      : Ptr64 _FILE_OBJECT
...

```

Iopb field, containing the operation's parameters (i.e. share access for IRP_MJ_CREATE or Length for IRP_MJ_READ) and the RequestorMode field, which tells us if the IRP was issued by the kernel or not.

The main difference between filtering I/O requests issued by win32 processes and Linux processes is that, for Linux processes, the RequestorMode is KernelMode, because IRPs are issued by the kernel, not by the actual process. That is because linux processes don't know of Windows IRPs, therefore the pico provider has to translate the I/O request into an IRP. Considering this, we filtered kernel issued IRPs too, but only those for which the requestor id was not system's PID. To do this, in the pre callback for every irp we checked PFLT_CALLBACK_DATA RequestorMode field, and if the requestor pid returned by FltGetRequestorProcessId is 4 (system process pid) or if the pid was of a win32 process that is not part of linux process tree, then we return FLT_PREOP_SUCCESS_NO_CALLBACK, thus skipping an unnecessary call to the post callback.

2) *Process Filtering*: In order to receive process creation and termination notifications, the driver must register a callback with PsSetCreateProcessNotifyRoutineEx2. This callback will be called synchronously for both win32 processes and pico processes. The algorithm on process creation is as follows:

- get the process' initial security token and check for initial privilege escalation
- determine if it is a WSL process
- insert the process in the process collector if it is a WSL process or if parent is in the collector
- insert the process in its parent process tree if it exists, else create a new process tree

Before any other processing is done, we need to check for a potential privilege escalation at process startup. We do this by checking if the process was elevated in any other way except by svchost or by an already elevated parent. Process elevation will be detailed in the privilege escalation detection section.

In order to identify a WSL process, the driver must call ZwQueryInformationProcess with SubsystemInformationTypeWSL information type on the process' handle and check that the subsystem type is SubsystemInformationTypeWSL (linux process) or SubsystemInformationTypeWin32 (win32 process). After determining the subsystem we insert an

entry into our process collector if the subsystem type is SubsystemInformationTypeWSL or if it's parent is already in the collector. This way we are sure that we monitor win32 processes too, if and only if they are part of a linux process tree. If its parent is already in the collector, we add it as a child process in its tree, else we create a new process tree.

VI. DETECTION HEURISTICS

A. Ransomware

Ransomware is a type of malware that encrypts the users' files and demands money, usually in the form of crypto currency, like bitcoin, in order for the user to regain access to the files. More advanced ransomware could go as far as preventing the operating system to boot until the payment is made. An example of such ransomware would be petya.

A ransomware that targets WSL would be leverage the fact that most AV minifilters do not filter IRPs that are coming from the kernel, allowing the ransomware to silently encrypt files. Even more, the ransomware might silently communicate with a server without being detected by network filters.

A reliable solution would be to monitor kernel issued IRPs that are done on behalf of a Linux process, and to use either a scoring engine or some AI expert system (i.e. Support Vector Machine) to identify the ransomware. We are going to cover only the behavioral heuristics approach in this paper.

In case of a scoring engine, the process could be assigned points according to the file system activity it does. For example, a file delete or write would add more points than a file read. Moreover, even more points could be added if the action involves sensitive paths (i.e. system root, Windows directory, etc). When the process reaches a certain threshold, a detection alert should be issued and the process should be killed. In order to detect multi-process ransomware we keep track of its whole process tree.

In order to do this, we keep a dictionary (internally, a trie) with all paths we monitor. Each path has an associated score for each I/O operation. The operations we are mostly interested in are file move, extension change, file delete, file write, file create. The score we keep per path per operation is heuristically chosen and can be fine tuned using different machine learning algorithms. Whenever an operation is filtered, we try to match the path against our dictionary, and add the corresponding points if the path was matched.

Matching paths for each file operation can be a costly operation. We had to come with a caching mechanism so that when we matched a path, we would know the scores we need to assign for any operation without rematching the path. We've decided that the best solution was using file contexts. The minifilter driver defines the context's structure and assigns it to a filter manager object (i.e. volume, file,

stream handle, etc). In our case, the context is a structure that contains the array of scores for the monitored operations. This way, whenever we filter a file operation, we check if it has a context associated with it (FltGetFileContext). If a context already exists, then we just add points for that process. If not, we allocate (FltAllocateContext) and set the context (FltSetFileContext), then add the points.

B. Windows Privilege Escalation

Privilege escalation is a powerful exploit that is meant to elevate an unelevated process without users' consent or knowledge. It is very important to understand the elevation process in order to identify the exploit. In this section we are going to describe in detail how elevation is done in Windows and how we can detect the exploit.

On Windows, there are multiple ways of starting an elevated process:

- from an unprivileged process, via ShellExecuteEx
- auto elevate via manifest
- from a privileged process

When an unprivileged process needs to start a privileged process, it sends a request to svchost, which shows the User Account Control pop-up through consent.exe, notifying the user that the process needs admin rights. If admin rights are granted by UAC, svchost starts the process. If a process already is elevated, any process spawned by it will also be elevated. We can easily see that, if wsl.exe is started with admin rights, all Linux processes running in that WSL instance will be granted admin rights inside Windows, which would be disastrous considering the security of the machine.

This kind of exploit can be detected easily, while not very reliable, from a kernel driver, or, more reliable, by using memory introspection techniques from a hypervisor. While the latter is extremely reliable, it is much more complex than a kernel-driver. We will cover only the first technique and try to make it as reliable as possible without adding too much performance overhead.

We need to check that the process' token was not tampered with. The algorithm for checking if a process is elevated is fairly simple. We use ZwOpenProcessTokenEx to get the process' security token and then ZwQueryInformationToken with TokenElevation information class. The returned boolean will tell if the process is elevated or not. This check can easily be done at in the process notification callback because we get the process' handle as a parameter.

Seeing how execve exploit works, it is clearly obvious that while this method works for the most basic case, it is quite unreliable. That's because we need to poll the security token at certain moments since the exploit just overwrites the process' security token with the System security token. Such key moments would be whenever we intercept actions that would naturally require admin privileges, i.e. accessing the

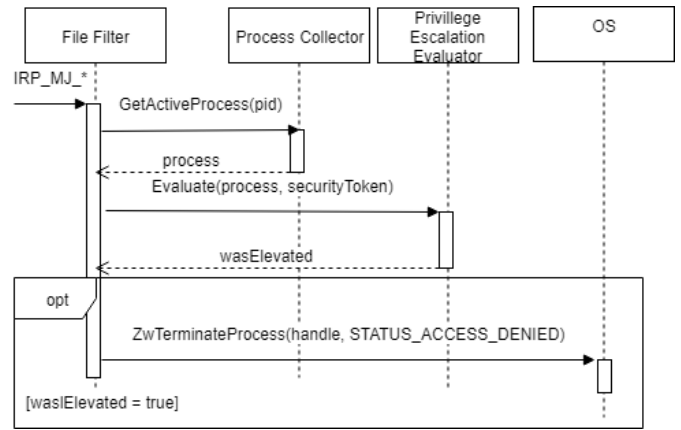


Fig. 2. privilege escalation detection diagram

Windows or System32 folders. It's easy to see several issues:

- if the exploit elevates the current process while running
- if the exploit is used to elevate another process
- if the elevated process is a Windows process
- performance overhead added by polling the security token

Further, we are going to explain how we've tackled these problems and made the algorithm more reliable.

1) *Detecting privilege escalation on already running processes:* In order to do this we need to keep track of all active processes in a structure that we'll call a process collector. When we add a new process into the collector, we will also store its initial elevation. Then, whenever the processes does sensitive actions, like dropping executables in paths like "C:\Windows", we will recheck the security token with the algorithm previously presented. If the new elevation differs from the initial elevation, the process should be detected and killed.

2) *Detecting elevation of a running Windows process:* We applied the same algorithm for Windows processes that are started by Linux processes as well, with two observations: We need to keep track of Windows processes as well, but only if they are part of a process tree that has a Linux process as well, and we need to detect and kill the whole process tree.

3) *Minimizing the performance overhead:* Polling the security token for every Linux process would become costly performance-wise very quickly, if the process does I/O intensive operations(i.e, a compiler). It is necessary to narrow the set of actions for which we do security token registration, as well as use an efficient algorithm for matching strings against a predefined dictionary.

C. File Infector Detection

A WSL file infector would essentially inject into windows executable files, altering the code in order to deliver some malicious code. Even though this can generally be prevented by software developers by signing their binaries, assuring that, if loaded, were not tampered with, unsigned binaries still exist and are vulnerable to being hijacked.

REFERENCES

- [1] Pavel Yosifovich, Alex Ionescu, Mark E. Russinovich, and David A. Solomon "Windows Internals"
- [2] Alex Ionescu "THE LINUX KERNEL HIDDEN INSIDE Windows 10," BlackHat2016
- [3] Alex Ionescu "GAINING VISIBILITY INTO LINUX BINARIES ON Windows: DEFEND AND UNDERSTAND WSL," BlueHat2016
- [4] Saar Amar "LINUX VULNERABILITIES, Windows EXPLOITS Escalating Privileges with WSL," BlueHat2018
- [5] Manuel Egele, Theodoor Scholte, Engin Kirda, and Christopher Kruegel "A Survey on Automated Dynamic Malware Analysis Techniques and Tools," ACM Computing Surveys (CSUR) ,February 2012